

▼ Informer Demo

▼ Download code and dataset

```
!git clone https://github.com/zhouhaoyi/Informer2020.git
!git clone https://github.com/zhouhaoyi/ETDataset.git
!ls
```

```
fatal: destination path 'Informer2020' already exists and is not an empty directory.
fatal: destination path 'ETDataset' already exists and is not an empty directory.
ETDataset  Informer2020  informer_checkpoints  results  sample_data
```

```
import sys
if not 'Informer2020' in sys.path:
    sys.path += ['Informer2020']
```

```
# !pip install -r ./Informer2020/requirements.txt
```

▼ Experiments: Train and Test

```
from utils.tools import dotdict
from exp.exp_informer import Exp_Informer
import torch
```

```
args = dotdict()
```

```
args.model = 'informer' # model of experiment, options: [informer, informerstack, informerlight(TBD)]
```

```
args.data = 'ETTh1' # data
args.root_path = './ETDataset/ETT-small/' # root path of data file
args.data_path = 'ETTh2.csv' # data file
args.features = 'S' # forecasting task, options:[M, S, MS]; M:multivariate predict multivariate, S:univariate predict univariate, MS:
args.target = 'OT' # target feature in S or MS task
args.freq = 'h' # freq for time features encoding, options:[s:secondly, t:minutely, h:hourly, d:daily, b:business days, w:weekly, m:m
args.checkpoints = './informer_checkpoints' # location of model checkpoints

args.seq_len = 336 # input sequence length of Informer encoder
args.label_len = 336 # start token length of Informer decoder
args.pred_len = 168 # prediction sequence length
# Informer decoder input: concat[start token series(label_len), zero padding series(pred_len)]

args.enc_in = 7 # encoder input size
args.dec_in = 7 # decoder input size
args.c_out = 7 # output size
args.factor = 5 # probsparse attn factor
args.d_model = 512 # dimension of model
args.n_heads = 8 # num of heads
args.e_layers = 2 # num of encoder layers
args.d_layers = 1 # num of decoder layers
args.d_ff = 1024 # dimension of fcn in model
args.dropout = 0.05 # dropout
args.attn = 'prob' # attention used in encoder, options:[prob, full]
args.embed = 'timeF' # time features encoding, options:[timeF, fixed, learned]
args.activation = 'gelu' # activation
args.distil = True # whether to use distilling in encoder
args.output_attention = False # whether to output attention in ecoder

args.batch_size = 32
args.learning_rate = 0.0001
args.loss = 'mse'
args.lradj = 'type1'
args.use_amp = False # whether to use automatic mixed precision training

args.num_workers = 0
args.itr = 1
```

```
args.train_epochs = 5
args.patience = 3
args.des = 'exp'

args.use_gpu = True if torch.cuda.is_available() else False
args.gpu = 0

args.use_multi_gpu = False
args.devices = '0,1,2,3'

args.use_gpu = True if torch.cuda.is_available() and args.use_gpu else False

if args.use_gpu and args.use_multi_gpu:
    args.dvices = args.devices.replace(' ','')
    device_ids = args.devices.split(',')
    args.device_ids = [int(id_) for id_ in device_ids]
    args.gpu = args.device_ids[0]

# Set augments by using data name
data_parser = {
    'ETTh1':{'data':'ETTh1.csv','T':'OT','M':[7,7,7],'S':[1,1,1],'MS':[7,7,1]},
    'ETTh2':{'data':'ETTh2.csv','T':'OT','M':[7,7,7],'S':[1,1,1],'MS':[7,7,1]},
    'ETTm1':{'data':'ETTm1.csv','T':'OT','M':[7,7,7],'S':[1,1,1],'MS':[7,7,1]},
    'ETTm2':{'data':'ETTm2.csv','T':'OT','M':[7,7,7],'S':[1,1,1],'MS':[7,7,1]},
}
if args.data in data_parser.keys():
    data_info = data_parser[args.data]
    args.data_path = data_info['data']
    args.target = data_info['T']
    args.enc_in, args.dec_in, args.c_out = data_info[args.features]

args.detail_freq = args.freq
args.freq = args.freq[-1:]
```


▼ Prediction

```
import os

# set saved model path
setting = 'informer_ETTh1_ftS_sl336_ll336_pl168_dm512_nh8_el2_dl1_df1024_atprob_fc5_ebtimeF_dtTrue_exp_0'
# path = os.path.join(args.checkpoints,setting,'checkpoint.pth')

# If you already have a trained model, you can set the arguments and model path, then initialize a Experiment and use it to predict
# Prediction is a sequence which is adjacent to the last date of the data, and does not exist in the data
# If you want to get more information about prediction, you can refer to code `exp/exp_informer.py` function predict() and `data/data

exp = Exp(args)

exp.predict(setting, True)

    Use GPU: cuda:0
    pred 1

# the prediction will be saved in ./results/{setting}/real_prediction.npy
import numpy as np

prediction = np.load('./results/'+setting+'/real_prediction.npy')

prediction.shape

(1, 168, 1)
```

▼ More details about Prediction - prediction function

```
# here is the detailed code of function predict
```

```

# HERE IS THE DETAILED CODE OF FUNCTION PREDICT

def predict(exp, setting, load=False):
    pred_data, pred_loader = exp._get_data(flag='pred')

    if load:
        path = os.path.join(exp.args.checkpoints, setting)
        best_model_path = path+'/'+ 'checkpoint.pth'
        exp.model.load_state_dict(torch.load(best_model_path))

    exp.model.eval()

    preds = []

    for i, (batch_x, batch_y, batch_x_mark, batch_y_mark) in enumerate(pred_loader):
        batch_x = batch_x.float().to(exp.device)
        batch_y = batch_y.float()
        batch_x_mark = batch_x_mark.float().to(exp.device)
        batch_y_mark = batch_y_mark.float().to(exp.device)

        # decoder input
        dec_inp = torch.zeros_like(batch_y[:, -exp.args.pred_len:, :]).float()
        dec_inp = torch.cat([batch_y[:, :, :exp.args.label_len, :], dec_inp], dim=1).float().to(exp.device)
        # encoder - decoder
        if exp.args.use_amp:
            with torch.cuda.amp.autocast():
                if exp.args.output_attention:
                    outputs = exp.model(batch_x, batch_x_mark, dec_inp, batch_y_mark)[0]
                else:
                    outputs = exp.model(batch_x, batch_x_mark, dec_inp, batch_y_mark)
        else:
            if exp.args.output_attention:
                outputs = exp.model(batch_x, batch_x_mark, dec_inp, batch_y_mark)[0]
            else:
                outputs = exp.model(batch_x, batch_x_mark, dec_inp, batch_y_mark)
        f_dim = -1 if exp.args.features=='MS' else 0
        batch_y = batch_y[:, -exp.args.pred_len:, f_dim:].to(exp.device)

```

```
pred = outputs.detach().cpu().numpy().squeeze()

preds.append(pred)

preds = np.array(preds)
preds = preds.reshape(-1, preds.shape[-2], preds.shape[-1])

# result save
folder_path = './results/' + setting + '/'
if not os.path.exists(folder_path):
    os.makedirs(folder_path)

np.save(folder_path+'real_prediction.npy', preds)


return preds

# you can also use this prediction function to get result
prediction = predict(exp, setting, True)

pred 1

import matplotlib.pyplot as plt

plt.figure()
plt.plot(prediction[0, :, -1])
plt.show()
```

▼ More details about Prediction - prediction dataset

You can give a `root_path` and `data_path` of the data you want to forecast, and set `seq_len`, `label_len`, `pred_len` and other arguments as other Dataset. The difference is that you can set a more detailed freq such as 15min or 3h to generate the timestamp of prediction series.

Dataset_Pred only has one sample (including `encoder_input: [1, seq_len, dim]`, `decoder_token: [1, label_len, dim]`, `encoder_input_timestamp: [1, seq_len, date_dim]`, `decoder_input_timestamp: [1, label_len+pred_len, date_dim]`). It will intercept the last sequence of the given data (`seq_len` data) to forecast the unseen future sequence (`pred_len` data).

```
from data.data_loader import Dataset_Pred
from torch.utils.data import DataLoader
```

```
Data = Dataset_Pred
timeenc = 0 if args.embed!='timeF' else 1
flag = 'pred'; shuffle_flag = False; drop_last = False; batch_size = 1
```

```
freq = args.detail_freq
```

```
data_set = Data(
    root_path=args.root_path,
    data_path=args.data_path,
    flag=flag,
    size=[args.seq_len, args.label_len, args.pred_len],
    features=args.features,
    target=args.target,
    timeenc=timeenc,
    freq=freq
)
data_loader = DataLoader(
```

```
data_set,  
batch_size=batch_size,  
shuffle=shuffle_flag,  
num_workers=args.num_workers,  
drop_last=drop_last)
```

```
len(data_set), len(data_loader)
```

```
(1, 1)
```

▼ Visualization

```
# When we finished exp.train(setting) and exp.test(setting), we will get a trained model and the results of test experiment  
# The results of test experiment will be saved in ./results/{setting}/pred.npy (prediction of test dataset) and ./results/{setting}/t
```

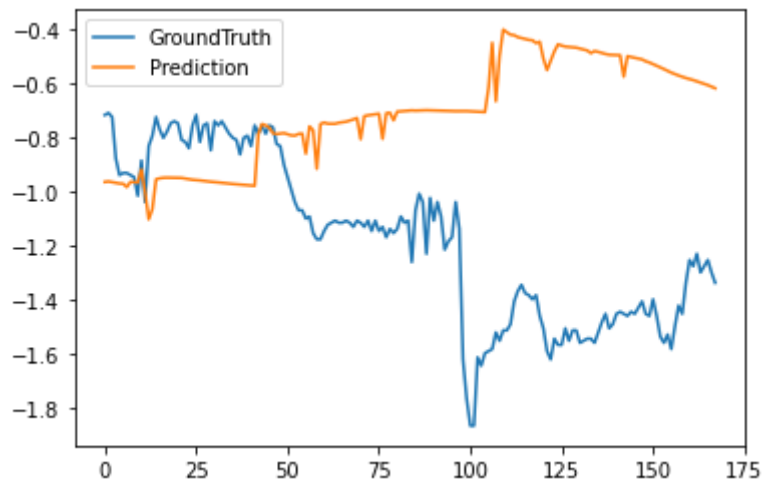
```
preds = np.load('./results/'+setting+'/pred.npy')  
trues = np.load('./results/'+setting+'/true.npy')
```

```
# [samples, pred_len, dimensions]  
preds.shape, trues.shape
```


```
((2688, 168, 1), (2688, 168, 1))
```

```
import matplotlib.pyplot as plt  
import seaborn as sns
```

```
# draw OT prediction  
plt.figure()  
plt.plot(trues[500,:,-1], label='GroundTruth')  
plt.plot(preds[500,:,-1], label='Prediction')  
plt.legend()  
plt.show()
```



```
# draw HUFL prediction
plt.figure()
plt.plot(trues[0,:,0], label='GroundTruth')
plt.plot(preds[0,:,0], label='Prediction')
plt.legend()
plt.show()
```



```
from data.data_loader import Dataset_ETT_hour
from torch.utils.data import DataLoader
```

```
Data = Dataset_ETT_hour
timeenc = 0 if args.embed!='timeF' else 1
flag = 'test'; shuffle_flag = False; drop_last = True; batch_size = 1

data_set = Data(
    root_path=args.root_path,
    data_path=args.data_path,
    flag=flag,
    size=[args.seq_len, args.label_len, args.pred_len],
    features=args.features,
    timeenc=timeenc,
    freq=args.freq
)
data_loader = DataLoader(
    data_set,
    batch_size=batch_size,
    shuffle=shuffle_flag,
    num_workers=args.num_workers,
    drop_last=drop_last)

import os

args.output_attention = True

exp = Exp(args)

model = exp.model

setting = 'informer_ETTh1_ftS_sl336_ll336_pl168_dm512_nh8_el2_dl1_df1024_atprob_fc5_ebtimeF_dtTrue_exp_0'
path = os.path.join(args.checkpoints,setting,'checkpoint.pth')
model.load_state_dict(torch.load(path))

    Use GPU: cuda:0
    <All keys matched successfully>
```

```
# attention visualization
idx = 0
for i, (batch_x, batch_y, batch_x_mark, batch_y_mark) in enumerate(data_loader):
    if i != idx:
        continue
    batch_x = batch_x.float().to(exp.device)
    batch_y = batch_y.float()

    batch_x_mark = batch_x_mark.float().to(exp.device)
    batch_y_mark = batch_y_mark.float().to(exp.device)

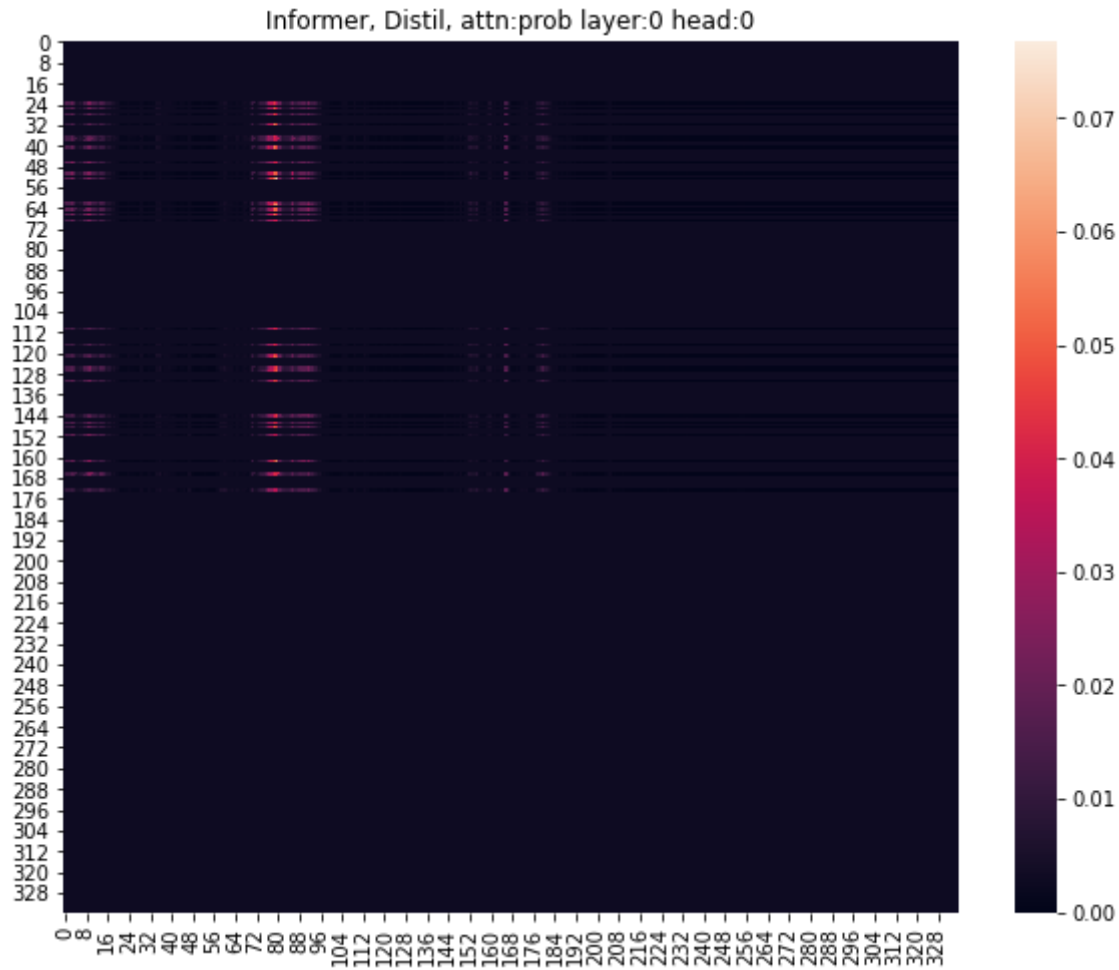
    dec_inp = torch.zeros_like(batch_y[:, -args.pred_len:, :]).float()
    dec_inp = torch.cat([batch_y[:, :, args.label_len:], dec_inp], dim=1).float().to(exp.device)

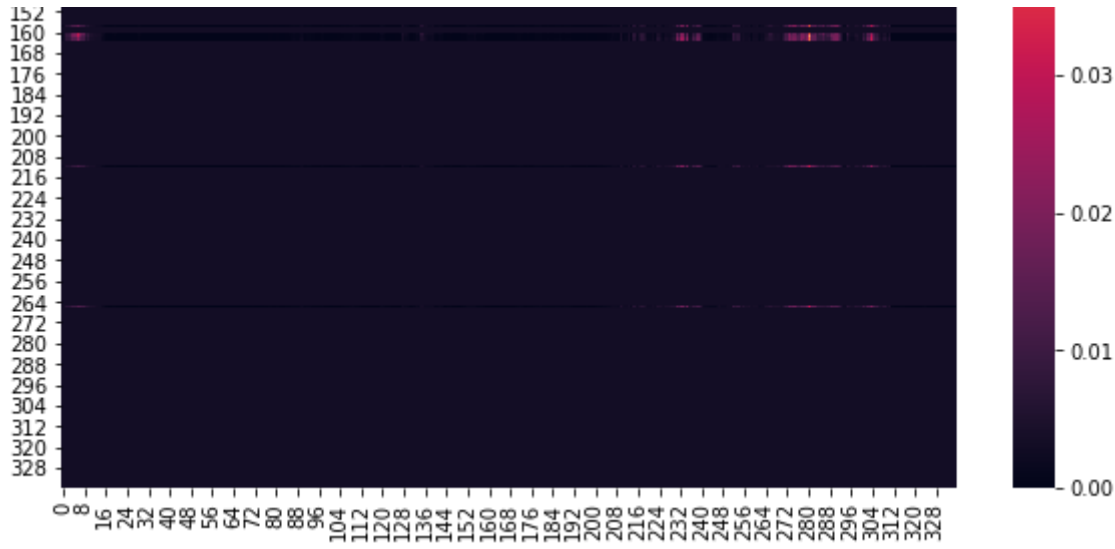
    outputs, attn = model(batch_x, batch_x_mark, dec_inp, batch_y_mark)

    attn[0].shape, attn[1].shape #, attn[2].shape

    (torch.Size([1, 8, 336, 336]), torch.Size([1, 8, 169, 169]))


    layer = 0
    distil = 'Distil' if args.distil else 'NoDistil'
    for h in range(0, 8):
        plt.figure(figsize=[10, 8])
        plt.title('Informer, {}, attn: {} layer: {} head: {}'.format(distil, args.attn, layer, h))
        A = attn[layer][0, h].detach().cpu().numpy()
        ax = sns.heatmap(A, vmin=0, vmax=A.max()+0.01)
        plt.show()
```



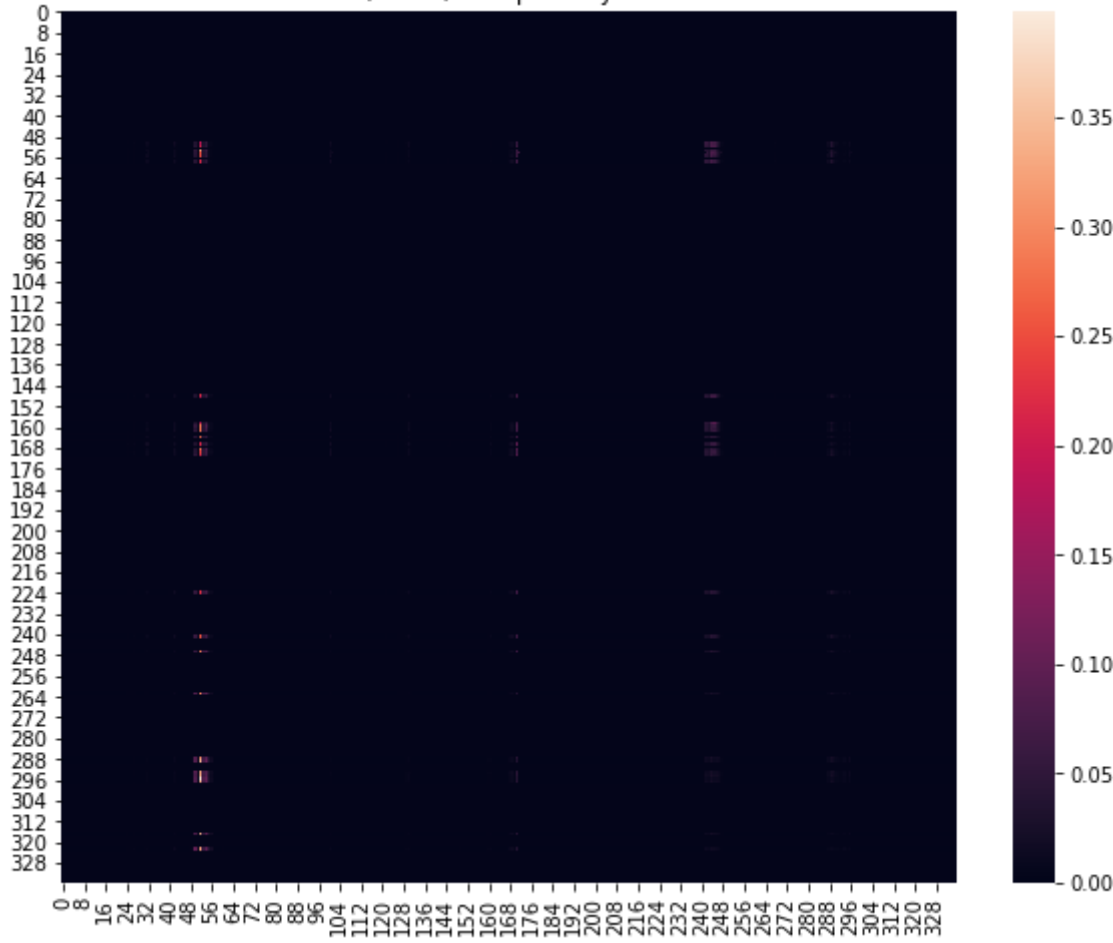


Informer, Distil, attn:prob layer:0 head:2







Informer, Distil, attn:prob layer:0 head:3





Informer, Distil, attn:prob layer:0 head:4






Informer, Distil, attn:prob layer:0 head:5





Informer, Distil, attn:prob layer:0 head:6



Informer, Distil, attn:prob layer:0 head:7





```
layer = 1
distil = 'Distil' if args.distil else 'NoDistil'
for h in range(0,8):
    plt.figure(figsize=[10,8])
    plt.title('Informer, {}, attn:{}'.format(distil, args.attn, layer, h))
    A = attn[layer][0,h].detach().cpu().numpy()
    ax = sns.heatmap(A, vmin=0, vmax=A.max()+0.01)
    plt.show()
```